

# gneiting\_-\_probabilistic\_forecasts\_calibration\_sharpness

Evan L. Ray

8/21/2020

## Article

This is about the article Probabilistic forecasts, calibration and sharpness by Gneiting et al. <https://rss.onlinelibrary.wiley.com/doi/10.1111/j.1467-9868.2007.00587.x>

## Summary

## Introduction

$G_t$  is data generating process,  $F_t$  is a forecasting model.

PIT is  $p_t = F_t(x_t)$ .  $p_t \sim Unif(0, 1)$  is a necessary but not sufficient condition for a forecaster to be ideal.

Example:

DGP is:

$$\begin{aligned}\mu_t &\sim N(0, 1) \\ X_t &\sim N(\mu_t, 1)\end{aligned}$$

Three other models in table.

```
r_dgp <- function(n) {  
  mu_t <- rnorm(n)  
  x_t <- rnorm(n = length(mu_t), mean = mu_t)  
  return(list(mu_t = mu_t, x_t = x_t))  
}  
  
d_dgp <- function(x_t, mu_t) {  
  dnorm(x = x_t, mean = mu_t, sd = 1)  
}  
  
p_dgp <- function(x_t, mu_t) {  
  if(missing(mu_t)) {  
    mu_t <- rnorm(length(x_t))  
  }  
  pnorm(q = x_t, mean = mu_t, sd = 1)  
}  
  
q_dgp <- function(p, mu_t) {  
  qnorm(p = p, mean = mu_t, sd = 1)  
}
```

```

}

r_f1 <- function(n) {
  list(
    x_t = rnorm(n = n, mean = 0, sd = sqrt(2))
  )
}

d_f1 <- function(x_t, ...) {
  dnorm(x = x_t, mean = 0, sd = sqrt(2))
}

p_f1 <- function(x_t, ...) {
  pnorm(q = x_t, mean = 0, sd = sqrt(2))
}

q_f1 <- function(p, ...) {
  qnorm(p = p, mean = 0, sd = sqrt(2))
}

r_f2 <- function(n) {
  mu_t <- rnorm(n)
  tau_t <- sample(x = c(-1, 1), size = n, replace = TRUE)
  x_t <- 0.5 * (rnorm(n = n, mean = mu_t) + rnorm(n = n, mean = mu_t + tau_t))
  return(list(mu_t = mu_t, tau_t = tau_t, x_t = x_t))
}

d_f2 <- function(x_t, mu_t, tau_t) {
  x_t <- 0.5 * (dnorm(x = x_t, mean = mu_t) + dnorm(x = x_t, mean = mu_t + tau_t))
}

p_f2 <- function(x_t, mu_t, tau_t) {
  if(missing(tau_t)) {
    tau_t <- sample(x = c(-1, 1), size = length(x_t), replace = TRUE)
  }
  x_t <- 0.5 * (pnorm(q = x_t, mean = mu_t) + pnorm(q = x_t, mean = mu_t + tau_t))
}

q_f2 <- function(p, mu_t, tau_t) {
  if(missing(tau_t)) {
    tau_t <- sample(x = c(-1, 1), size = length(p), replace = TRUE)
  }
  x_t <- 0.5 * (qnorm(p = p, mean = mu_t) + pnorm(q = x_t, mean = mu_t + tau_t))
}

r_f3 <- function(n) {
  mu_t <- rnorm(n)
  mix_ind <- sample(x = 1:3, size = n, replace = TRUE)
  delta_t <- c(0.5, -0.5, 0)[mix_ind]
  sigma_sq_t <- c(1, 1, 169/100)[mix_ind]

  list(
    x_t = rnorm(n = n, mean = mu_t + delta_t, sd = sqrt(sigma_sq_t)),

```

```

    mix_ind = mix_ind
  )
}

d_f3 <- function(x_t, mu_t, mix_ind) {
  delta_t <- c(0.5, -0.5, 0)[mix_ind]
  sigma_sq_t <- c(1, 1, 169/100)[mix_ind]

  dnorm(x = x_t, mean = mu_t + delta_t, sd = sqrt(sigma_sq_t))
}

p_f3 <- function(x_t, mu_t, mix_ind) {
  if(missing(mix_ind)) {
    mix_ind <- sample(x = 1:3, size = length(x_t), replace = TRUE)
  }
  delta_t <- c(0.5, -0.5, 0)[mix_ind]
  sigma_sq_t <- c(1, 1, 169/100)[mix_ind]

  pnorm(q = x_t, mean = mu_t + delta_t, sd = sqrt(sigma_sq_t))
}

q_f3 <- function(p, mu_t, mix_ind) {
  delta_t <- c(0.5, -0.5, 0)[mix_ind]
  sigma_sq_t <- c(1, 1, 169/100)[mix_ind]

  qnorm(p = p, mean = mu_t + delta_t, sd = sqrt(sigma_sq_t))
}

```

Example predictive distributions for a single time point:

```

set.seed(42)
sample_dgp <- r_dgp(n = 1)
sample_f1 <- r_f1(n = 1)
sample_f2 <- r_f2(n = 1)
sample_f3 <- r_f3(n = 1)

p1 <- ggplot(data = data.frame(x = c(-5, 5)), mapping = aes(x = x)) +
  stat_function(
    fun = d_dgp, args = list(mu_t = sample_dgp$mu_t)
  ) +
  stat_function(
    fun = d_f1,
    color = "blue",
    linetype = 2
  ) +
  stat_function(
    fun = d_f2, args = list(mu_t = sample_dgp$mu_t, tau_t = sample_f2$tau_t),
    color = "cornflowerblue",
    linetype = 3
  ) +
  stat_function(

```


```

    fun = d_f3, args = list(mu_t = sample_dgp$mu_t, mix_ind = sample_f3$mix_ind),
    color = "orange",
    linetype = 4
  ) +
  ggtitle("pdfs") +
  theme_bw()

p2 <- ggplot(data = data.frame(x = c(-5, 5)), mapping = aes(x = x)) +
  stat_function(
    fun = p_dgp, args = list(mu_t = sample_dgp$mu_t)
  ) +
  stat_function(
    fun = p_f1,
    color = "blue",
    linetype = 2
  ) +
  stat_function(
    fun = p_f2, args = list(mu_t = sample_dgp$mu_t, tau_t = sample_f2$tau_t),
    color = "cornflowerblue",
    linetype = 3
  ) +
  stat_function(
    fun = p_f3, args = list(mu_t = sample_dgp$mu_t, mix_ind = sample_f3$mix_ind),
    color = "orange",
    linetype = 4
  ) +
  ggtitle("cdfs") +
  theme_bw()

gridExtra::grid.arrange(p1, p2)

```




Marginal predictive distributions:

```
n <- 100000

samples <- purrr::map_dfr(
  c("dgp", paste0("f", 1:3)),
  function(dist_name) {
    data.frame(
      dist = dist_name,
      x = do.call(
        paste0("r_", dist_name),
        args = list(n = n)
      )$x_t,
      stringsAsFactors = FALSE
    )
  }
)

ggplot(data = samples, mapping = aes(x = x, color = dist)) +
  geom_density() +
  theme_bw()
```



## PITs

```

samples <- r_dgp(n)

pits <- purrr::map_dfr(
  c("dgp", paste0("f", 1:3)),
  function(dist_name) {
    data.frame(
      dist = dist_name,
      pit = do.call(
        paste0("p_", dist_name),
        args = samples
      )
    )
  }
)

## Warning in bind_rows_(x, .id): Unequal factor levels: coercing to character
## Warning in bind_rows_(x, .id): binding character and factor vector, coercing
## into character vector

## Warning in bind_rows_(x, .id): binding character and factor vector, coercing
## into character vector

## Warning in bind_rows_(x, .id): binding character and factor vector, coercing


```

```
## into character vector
```

```
## Warning in bind_rows(x, .id): binding character and factor vector, coercing  
## into character vector
```

```
ggplot(data = pits, mapping = aes(x = pit)) +  
  geom_histogram(boundary = 0) +  
  facet_wrap( ~ dist)
```

```
## `stat_bin()` using `bins = 30`. Pick better value with `binwidth`.
```



PITs, but making the marginal model the DGP and the original DGP a predictive model

```
samples <- r_f1(n)  
dgp_samples <- r_dgp(n)  
  
pits <- purrr::map_dfr(  
  c("dgp", paste0("f", 1:3)),  
  function(dist_name) {  
    call_args <- samples  
    if(dist_name != "f1") {  
      call_args$mu_t <- dgp_samples$mu_t  
    }  
    data.frame(  
      dist = dist_name,  
      pit = do.call(  

```

```

    paste0("p_", dist_name),
    args = call_args
  )
}
)

```

```

## Warning in bind_rows_(x, .id): Unequal factor levels: coercing to character
## Warning in bind_rows_(x, .id): binding character and factor vector, coercing
## into character vector

## Warning in bind_rows_(x, .id): binding character and factor vector, coercing
## into character vector

## Warning in bind_rows_(x, .id): binding character and factor vector, coercing
## into character vector

## Warning in bind_rows_(x, .id): binding character and factor vector, coercing
## into character vector

```

```


ggplot(data = pits, mapping = aes(x = pit)) +
  geom_histogram(boundary = 0) +
  facet_wrap( ~ dist)

```

```

## `stat_bin()` using `bins = 30`. Pick better value with `binwidth`.

```





## Section 2: Modes of Calibration

climatological forecaster (marginal distribution) is probabilistically calibrated

We pick a particular value  $p = 0.75$  at which to explore probabilistic calibration.

```
prob_val <- 0.75

set.seed(42)

make_prob_cal_plot <- function() {
  sample_dgp <- r_dgp(n = 1)

  x_val <- q_f1(p = prob_val)

  if(x_val > 0) {
    dgp_lines <- data.frame(
      x = c(0, rep(x_val, 2)),
      y = c(rep(p_dgp(x_val, mu_t = sample_dgp$mu_t), 2), 0),
      model = "dgp"
    )

    f1_lines <- data.frame(
      x = c(0, rep(x_val, 2)),
      y = c(rep(prob_val, 2), 0)
    )
  } else {
    dgp_lines <- data.frame(
      x = c(rep(x_val, 2), 0),
      y = c(0, prob_val, prob_val),
      model = "dgp"
    )

    f1_lines <- data.frame(
      x = c(rep(x_val, 2), 0),
      y = c(0, rep(p_f1(x_val), 2))
    )
  }

  p <- ggplot(data = data.frame(x = c(-5, 5)), mapping = aes(x = x)) +
    geom_hline(yintercept = 0) +
    geom_vline(xintercept = 0) +
    stat_function(
      fun = p_dgp, args = list(mu_t = sample_dgp$mu_t),
      color = "orange"
    ) +
    stat_function(
      fun = p_f1,
      color = "blue",
      linetype = 2
    ) +
    geom_line(data = dgp_lines, mapping = aes(x = x, y = y),
              color = "orange") +
    geom_line(data = f1_lines, mapping = aes(x = x, y = y),
```

```

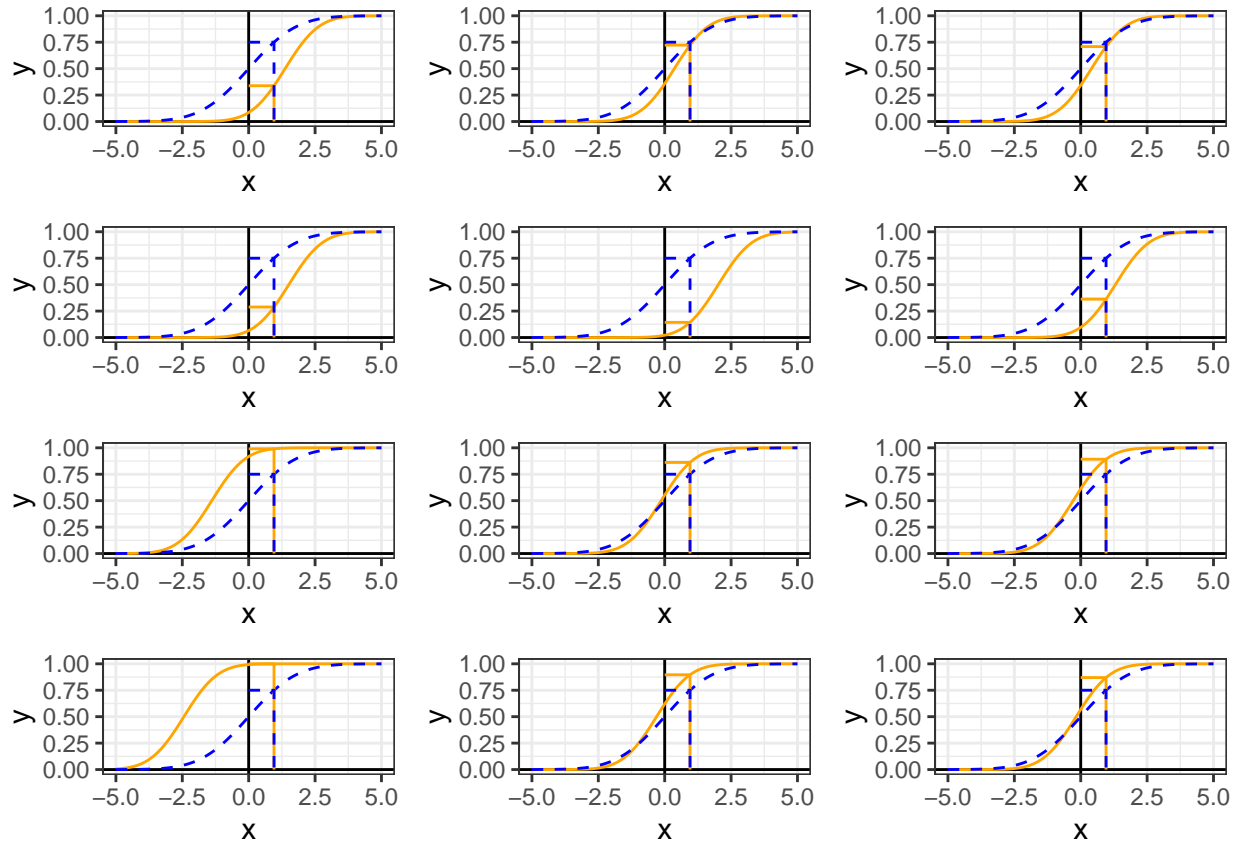
        color = "blue", linetype = 2) +
    theme_bw()

    return(p)
}

plots <- lapply(1:12, function(i) make_prob_cal_plot())

grid.arrange(grobs = plots)

```



Probabilistic calibration: at each probability level  $p$ , on average across all time points, the true probability of being less than the  $p$ 'th quantile of the predictive distribution is  $p$ . Note that this does not have to hold for every time point.

```

set.seed(42)
get_f1_prob_cal_rel_dgp <- function() {
  sample_dgp <- r_dgp(n = 1)

  x_val <- q_f1(p = prob_val)


  return(p_dgp(x_val, mu_t = sample_dgp$mu_t))
}

to_plot <- data.frame(
  p = purrr::map_dbl(1:100000, function(i) get_f1_prob_cal_rel_dgp())
)

```

```
ggplot(data = to_plot, mapping = aes(x = p)) +
  geom_histogram() +
  geom_vline(xintercept = mean(to_plot$p), color = "orange")
```

## `stat\_bin()` using `bins = 30`. Pick better value with `binwidth`.



Note: I accidentally code it up backwards at first and the relationship does not work in reverse!

```
set.seed(42)
get_f1_prob_cal_rel_dgp <- function() {
  sample_dgp <- r_dgp(n = 1)


  x_val <- q_dgp(p = prob_val, mu_t = sample_dgp$mu_t)

  return(p_f1(x_val))
}

to_plot <- data.frame(
  p = purrr::map_dbl(1:100000, function(i) get_f1_prob_cal_rel_dgp())
)

ggplot(data = to_plot, mapping = aes(x = p)) +
  geom_histogram() +
  geom_vline(xintercept = mean(to_plot$p), color = "orange")
```

## `stat\_bin()` using `bins = 30`. Pick better value with `binwidth`.



This is also an illustration that marginal calibration does not imply probabilistic calibration.

### Marginal distribution is exceedance calibrated

We pick a particular value  $x = 0.75$  at which to explore exceedance calibration.

```
x_val <- 0.75

set.seed(42)

make_prob_cal_plot <- function() {
  sample_dgp <- r_dgp(n = 1)

  x_val <- q_f1(p = prob_val)

  if(x_val > 0) {
    dgp_lines <- data.frame(
      x = c(0, rep(x_val, 2)),
      y = c(rep(p_dgp(x_val, mu_t = sample_dgp$mu_t), 2), 0),
      model = "dgp"
    )

    f1_lines <- data.frame(
      x = c(0, rep(x_val, 2)),
      y = c(rep(prob_val, 2), 0)
    )
  } else {
```

```

dgp_lines <- data.frame(
  x = c(rep(x_val, 2), 0),
  y = c(0, prob_val, prob_val),
  model = "dgp"
)

f1_lines <- data.frame(
  x = c(rep(x_val, 2), 0),
  y = c(0, rep(p_f1(x_val), 2))
)
}

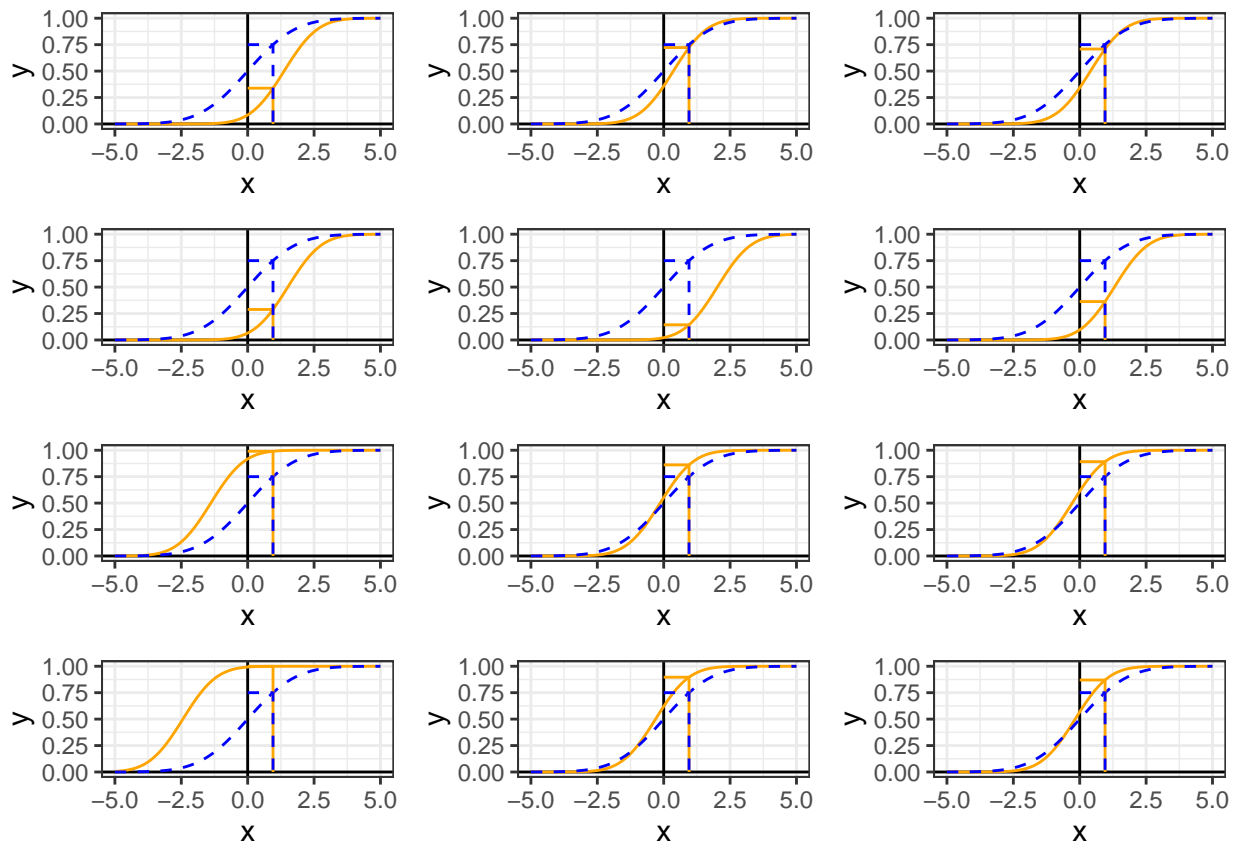
p <- ggplot(data = data.frame(x = c(-5, 5)), mapping = aes(x = x)) +
  geom_hline(yintercept = 0) +
  geom_vline(xintercept = 0) +
  stat_function(
    fun = p_dgp, args = list(mu_t = sample_dgp$mu_t),
    color = "orange"
  ) +
  stat_function(
    fun = p_f1,
    color = "blue",
    linetype = 2
  ) +
  geom_line(data = dgp_lines, mapping = aes(x = x, y = y),
    color = "orange") +
  geom_line(data = f1_lines, mapping = aes(x = x, y = y),
    color = "blue", linetype = 2) +
  theme_bw()

return(p)
}

plots <- lapply(1:12, function(i) make_prob_cal_plot())

grid.arrange(grobs = plots)

```



Probabilistic calibration: at each probability level  $p$ , on average across all time points, the true probability of being less than the  $p$ 'th quantile of the predictive distribution is  $p$ . Note that this does not have to hold for every time point.

```
set.seed(42)
get_f1_prob_cal_rel_dgp <- function() {
  sample_dgp <- r_dgp(n = 1)


  x_val <- q_f1(p = prob_val)

  return(p_dgp(x_val, mu_t = sample_dgp$mu_t))
}

to_plot <- data.frame(
  p = purrr::map_dbl(1:100000, function(i) get_f1_prob_cal_rel_dgp())
)

ggplot(data = to_plot, mapping = aes(x = p)) +
  geom_histogram() +
  geom_vline(xintercept = mean(to_plot$p), color = "orange")

## `stat_bin()` using `bins = 30`. Pick better value with `binwidth`.
```



Note: I accidentally code it up backwards at first and the relationship does not work in reverse!

```
set.seed(42)
get_f1_prob_cal_rel_dgp <- function() {
  sample_dgp <- r_dgp(n = 1)


  x_val <- q_dgp(p = prob_val, mu_t = sample_dgp$mu_t)

  return(p_f1(x_val))
}

to_plot <- data.frame(
  p = purrr::map_dbl(1:100000, function(i) get_f1_prob_cal_rel_dgp())
)

ggplot(data = to_plot, mapping = aes(x = p)) +
  geom_histogram() +
  geom_vline(xintercept = mean(to_plot$p), color = "orange")

## `stat_bin()` using `bins = 30`. Pick better value with `binwidth`.
```



This is also an illustration that marginal calibration does not imply probabilistic calibration.